

Implementing Page Rank Algorithm using MPI

Purshottam Vishwakarma

School of Informatics and Computing
Indiana University, Bloomington
pvishwak@indiana.edu

Bitan Saha

School of Informatics and Computing
Indiana University, Bloomington
bsaha@indiana.edu

Introduction:

Page Rank Algorithm, developed by Larry Page and Sergey Brin at Stanford University as part of their research, is the brain behind Google search engine. It uses probability distribution across all the pages over the web and assigns a value between 0 and 1 to each web page which determines the popularity of the page. Computation of page rank is an iterative process and the accuracy of the page rank is directly proportional to the number of iterations performed. Considering billions of web pages over the internet it is a highly compute intensive task and requires the algorithm to be run in parallel over a distributed environment. This Project implements the sequential version of the algorithm in a far smaller scale i.e. we would rank only 1000 web pages. *We have tried to make this algorithm highly efficient (3x faster) as compared to our peers.*

Theory:

In the general case, the PageRank value for any page i can be expressed as:

$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

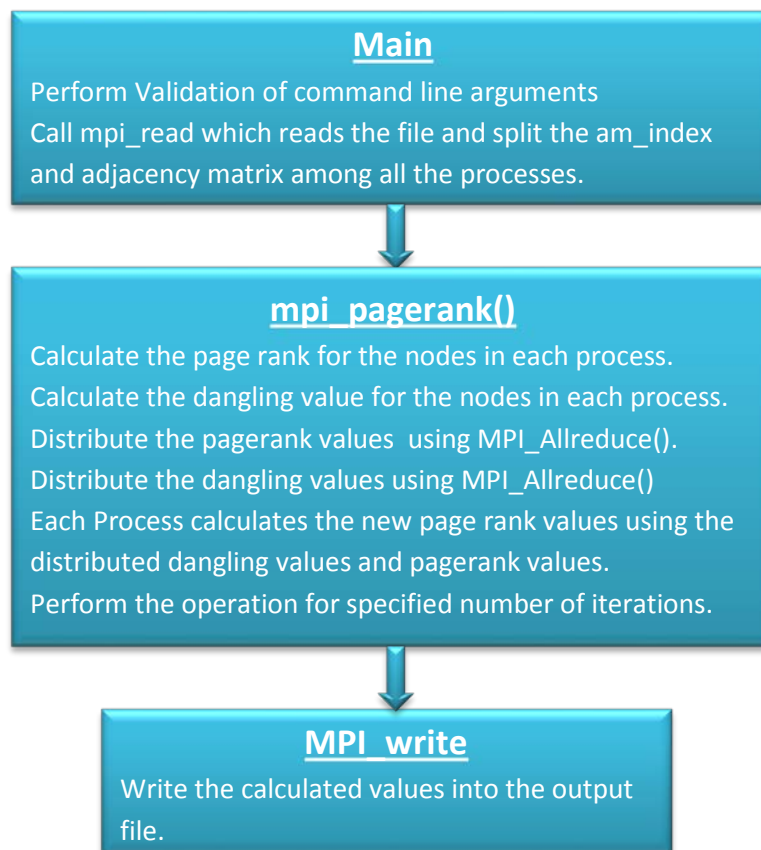
PR, pagerank (a probability value)

p_i , a page under consideration

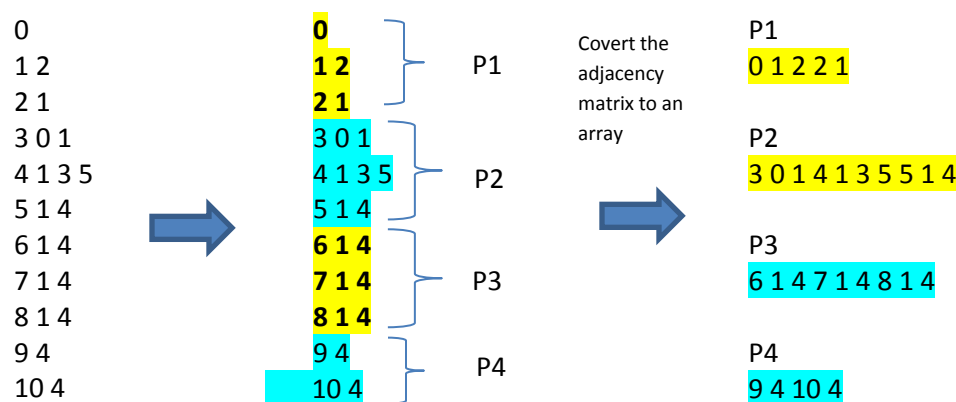
$L(p_i)$, the number of outbound links on page p_i

d , damping factor which can be set between 0 and 1 (It is usually set d to 0.85)

N , total number of pages.

Architecture design (e.g. architecture flowchart and text)**Implementation:**

- The algorithm splits the entire adjacency matrix among the number of processes specified.
- For example, the below adjacency matrix is split among 4 processes in the following manner.



- Each process keeps the track of the root node (1st column in the original adjacency matrix) using am_index data structure.
- Each process calculates the associate the page rank of the outbound nodes and dangling values of the root node. After calculating them each process sends these two values

(pagerank_values_table and the dangling values) to MPI and it sums up the values received from all the processes and then sends back the aggregated values to all the processes. Thus each process now has the aggregated values of the pagerank.

- Similarly MPI_Allreduce() receives individual process dangling values, sums them up and sends back the aggregate values back to each process.
- Each process then adds the damping factor into the calculated value to get the final value of page rank.
- The above steps are performed for the specified number of iterations.

Below is the diagrammatic representation of the algorithm for **iteration 1**.

| Process 1 | Process 2 | Process 3 | Process 4 |
|---------------------------|--------------------|--------------------|--------------------|
| [0] = 0.000000 | [0] = 0.045455 | [0] = 0.000000 | [0] = 0.000000 |
| [1] = 0.090909 | [1] = 0.121212 | [1] = 0.136364 | [1] = 0.000000 |
| [2] = 0.090909 | [2] = 0.000000 | [2] = 0.000000 | [2] = 0.000000 |
| [3] = 0.000000 | [3] = 0.030303 | [3] = 0.000000 | [3] = 0.000000 |
| [4] = 0.000000 | [4] = 0.045455 | [4] = 0.136364 | [4] = 0.181818 |
| [5] = 0.000000 | [5] = 0.030303 | [5] = 0.000000 | [5] = 0.000000 |
| [6] = 0.000000 | [6] = 0.000000 | [6] = 0.000000 | [6] = 0.000000 |
| [7] = 0.000000 | [7] = 0.000000 | [7] = 0.000000 | [7] = 0.000000 |
| [8] = 0.000000 | [8] = 0.000000 | [8] = 0.000000 | [8] = 0.000000 |
| [9] = 0.000000 | [9] = 0.000000 | [9] = 0.000000 | [9] = 0.000000 |
| [10] = 0.000000 | [10] = 0.000000 | [10] = 0.000000 | [10] = 0.000000 |
| Dangling value = 0.090909 | Dangling value = 0 | Dangling value = 0 | Dangling value = 0 |



MPI_Allreduce

| |
|-------------------------|
| [0] = 0.045455 |
| [1] = 0.348485 |
| [2] = 0.090909 |
| [3] = 0.030303 |
| [4] = 0.363636 |
| [5] = 0.030303 |
| [6] = 0.000000 |
| [7] = 0.000000 |
| [8] = 0.000000 |
| [9] = 0.000000 |
| [10] = 0.000000 |
| Sum_Dangling = 0.090909 |

Performance Analysis:

We have benchmarked our program with the sequential version of the program and found that the MPI version of page rank works faster than the sequential version. Below are the timings for the sequential version and the MPI version with 1000 nodes, 10 processes and 10 iterations.

| | |
|--|------|
| Sequential Version of PageRank | 1 s |
| MPI Version of PageRank (5 processes) | 1 ms |
| MPI Version of PageRank (10 processes) | 3 ms |
| MPI Version of PageRank (15 processes) | 4 ms |
| MPI Version of PageRank (20 processes) | 8 ms |

Results:

Below are the page ranks of the top 10 ranking URL numbers.

| | |
|-----|----------|
| 4 | 0.134060 |
| 34 | 0.125075 |
| 0 | 0.111998 |
| 20 | 0.084226 |
| 146 | 0.066848 |
| 2 | 0.048978 |
| 12 | 0.021937 |
| 14 | 0.017361 |
| 16 | 0.012771 |
| 66 | 0.012066 |

Acknowledgements:

We acknowledge the efforts of our Instructor **Judy Qiu** and Associate Instructors **Ikhyun Park & and Tak Lon Wu** in making us understand the basic concepts behind page rank.

References

1. Sergey Brin and Lawrence Page, [The Anatomy of a Large-Scale Hypertextual Web Search Engine](#), Stanford University, WWW7 Proceedings of the seventh international conference on World Wide Web 7, 1998
2. http://en.wikipedia.org/wiki/Markov_chain
3. http://en.wikipedia.org/wiki/Adjacency_matrix
4. <http://en.wikipedia.org/wiki/PageRank>
5. <http://www.open-mpi.org/doc/v1.4/>