

Assignment # 3 : Memory Hierarchy Optimizations

Report

- Deliverables :**

| | | |
|------------------|--|--|
| matrix_block.cpp | Matrix Multiplication (maximum 1100 X 1100) implementing blocking with different Block sizes and different matrix sizes. | |
| Datalayout.cpp | Sample program implementing two different data layouts. | |
| | Layout 1 | Layout 2 |
| | <pre>struct Sample1{ int x; double y[100]; }; Sample1 s[SIZE]; for(i=0;i<SIZE;i++) s[i].x = 25;</pre> | <pre>struct Sample2{ int x[SIZE]; double y[100]; }; Sample2 s2; for(j=0;j<SIZE;j++) s2.x[i] = 25;</pre> |
| | | |
| Makefile | Creates the executables matrix_block.out and datalayout.out | |

- NOTE:**

The programs were run on PUMBAA with below configuration:

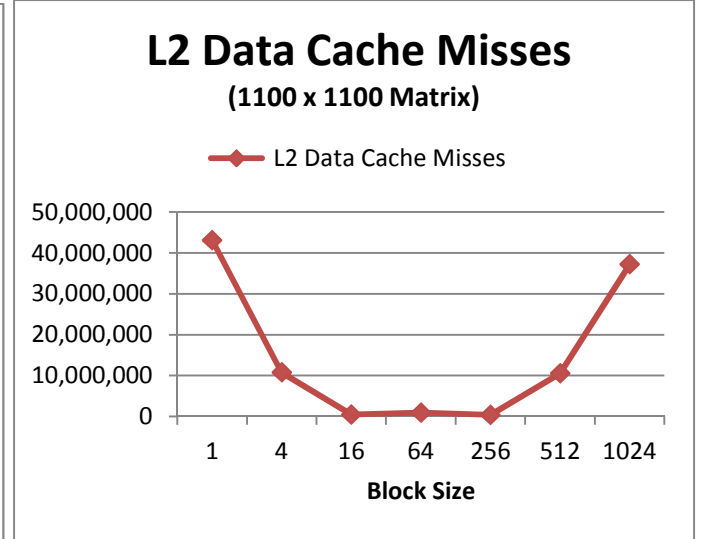
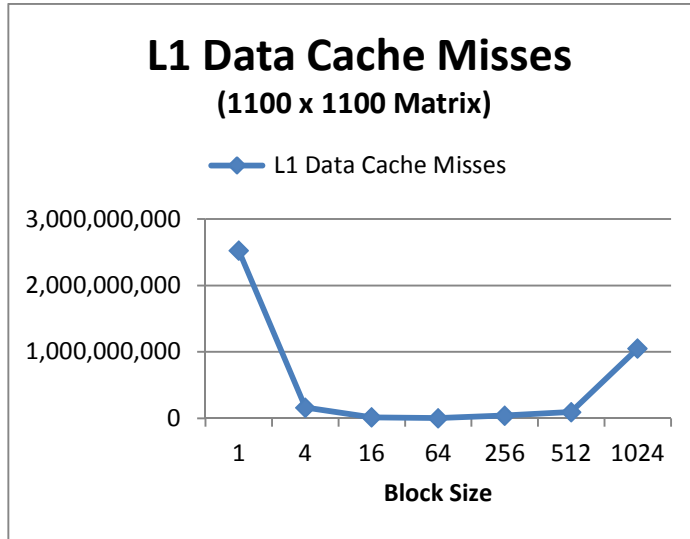
8 Core Machine:

Cache line size = 64 bytes on each index (Total 4 index on each core)

- Analysis for Part 1:**

Matrix size = 1100 x 1100

| Block Size | L1 Data Cache Misses | L2 Data Cache Misses |
|------------|----------------------|----------------------|
| 1 | 2524952177 | 43097142 |
| 4 | 160412344 | 10747937 |
| 16 | 14998730 | 454398 |
| 64 | 1727098 | 923593 |
| 256 | 41885974 | 373443 |
| 512 | 94004584 | 10526208 |
| 1024 | 1050323666 | 37207257 |



Explanation :

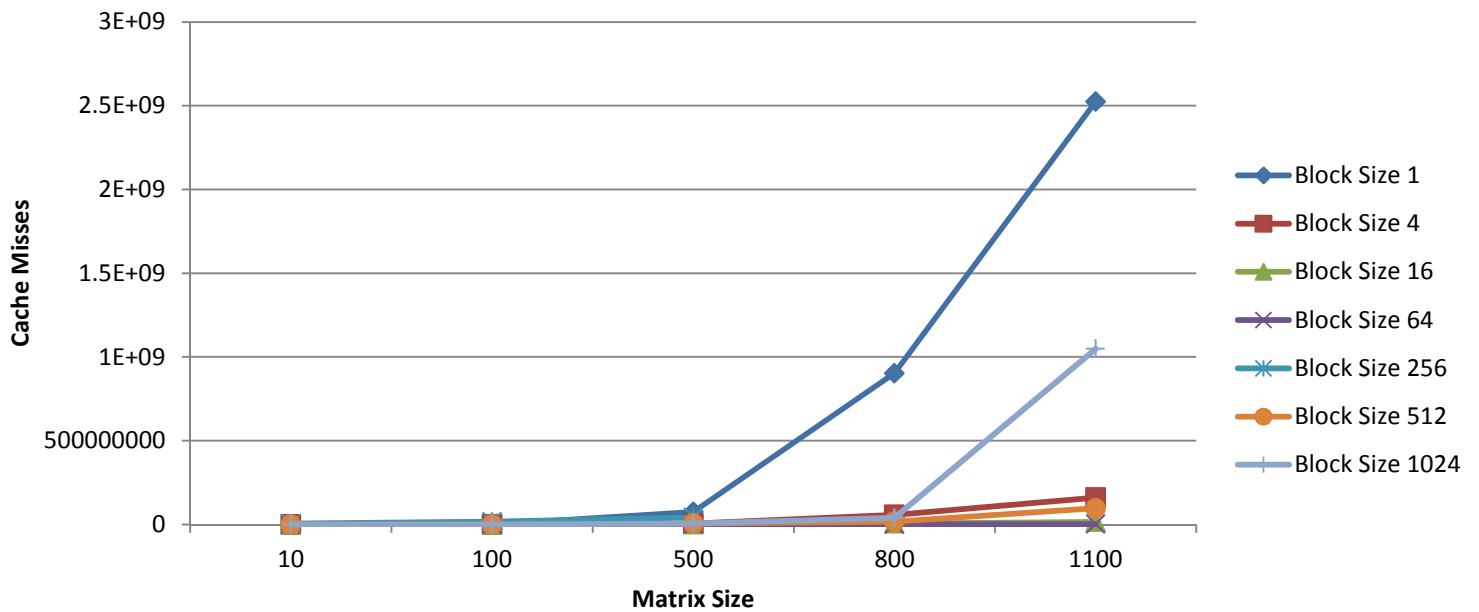
- As the block size increases from 1, subsequent data is available in cache due to temporal locality. Hence reduction in the number of cache misses.
- But as the block size increases beyond a certain level (in the above case 256) the amount of data that can be accommodated in a cache line is limited. Hence, rise in cache misses.

Analysis with Different Matrix Sizes:

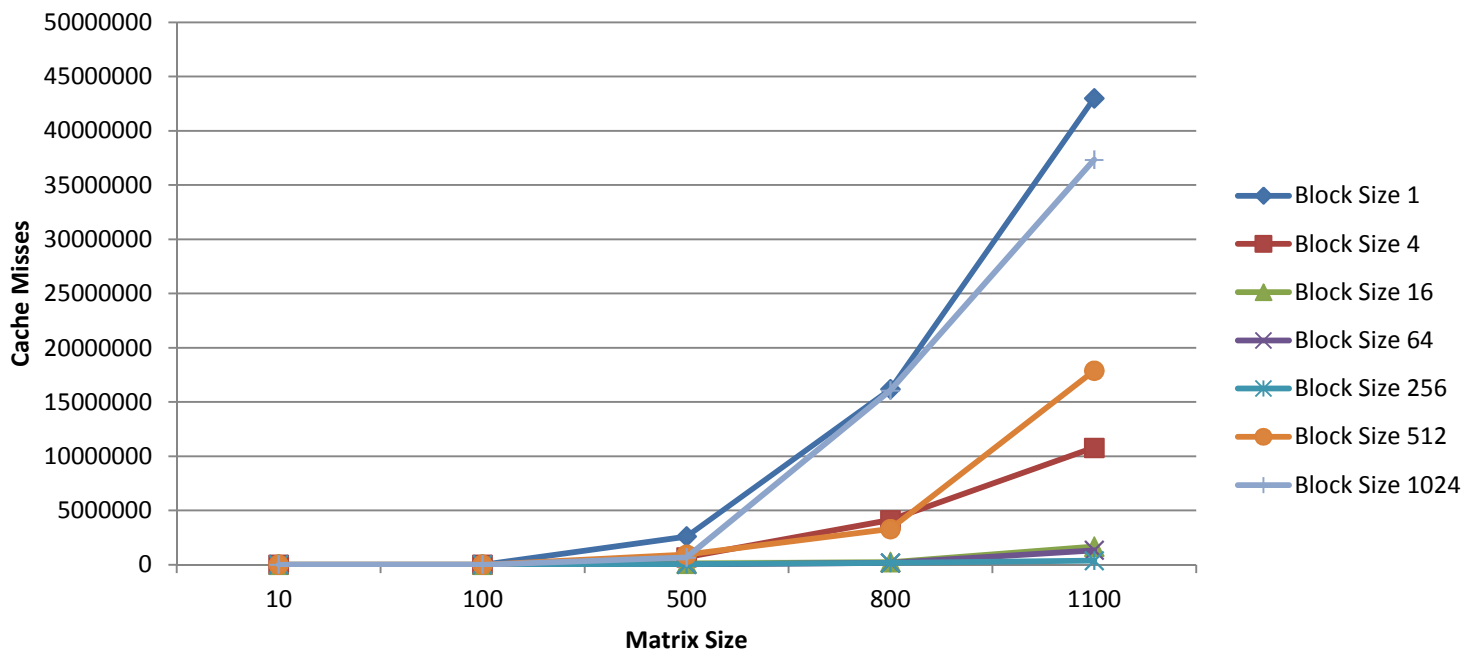
| L1 Data Cache Misses | | | | | | | |
|----------------------|--------------|--------------|---------------|---------------|----------------|----------------|-----------------|
| Matrix Size | Block Size 1 | Block Size 4 | Block Size 16 | Block Size 64 | Block Size 256 | Block Size 512 | Block Size 1024 |
| 10 | 85 | 91 | 96 | 84 | 134 | 73 | 98 |
| 100 | 1433 | 1344 | 1428 | 917 | 588 | 527 | 531 |
| 500 | 75039562 | 5630696 | 1308354 | 158001 | 4407889 | 5467973 | 5500649 |
| 800 | 902736743 | 57114256 | 6044457 | 933019 | 16616791 | 17637483 | 40194796 |
| 1100 | 2524934907 | 160360162 | 15202815 | 2190729 | 42550320 | 96241422 | 1049286914 |

| L2 Data Cache Misses | | | | | | | |
|----------------------|--------------|--------------|---------------|---------------|----------------|----------------|-----------------|
| Matrix Size | Block Size 1 | Block Size 4 | Block Size 16 | Block Size 64 | Block Size 256 | Block Size 512 | Block Size 1024 |
| 10 | 58 | 70 | 57 | 59 | 63 | 60 | 54 |
| 100 | 992 | 994 | 911 | 479 | 415 | 409 | 413 |
| 500 | 2593040 | 692807 | 103808 | 34903 | 40027 | 927808 | 670625 |
| 800 | 16168224 | 4109630 | 234950 | 168403 | 170307 | 3285606 | 16055812 |
| 1100 | 42986243 | 10784785 | 1671032 | 1328504 | 375480 | 17880677 | 37324339 |

L1 Data Cache Misses



L2 Data Cache Misses



• Analysis for Part 2:

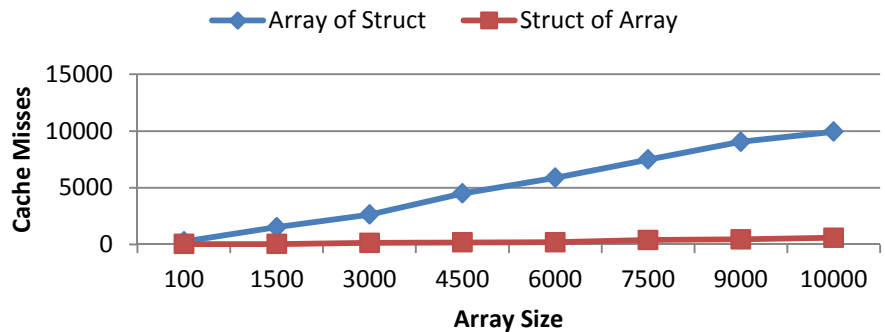
Sample program implementing two different data layouts.

| Layout 1 | Layout 2 |
|--|--|
| <pre>struct Sample1{ int x; double y[100]; }; Sample1 s[SIZE]; for(i=0;i<SIZE;i++) s[i].x = 25;</pre> | <pre>struct Sample2{ int x[SIZE]; double y[100]; }; Sample2 s2; for(j=0;j<SIZE;j++) s2.x[i] = 25;</pre> |

L1 Data Cache Misses

| Array Size | Array of Struct | Struct of Array |
|------------|-----------------|-----------------|
| 100 | 263 | 48 |
| 1500 | 1520 | 40 |
| 3000 | 2634 | 141 |
| 4500 | 4494 | 183 |
| 6000 | 5880 | 192 |
| 7500 | 7481 | 393 |
| 9000 | 9056 | 448 |
| 10000 | 9939 | 593 |

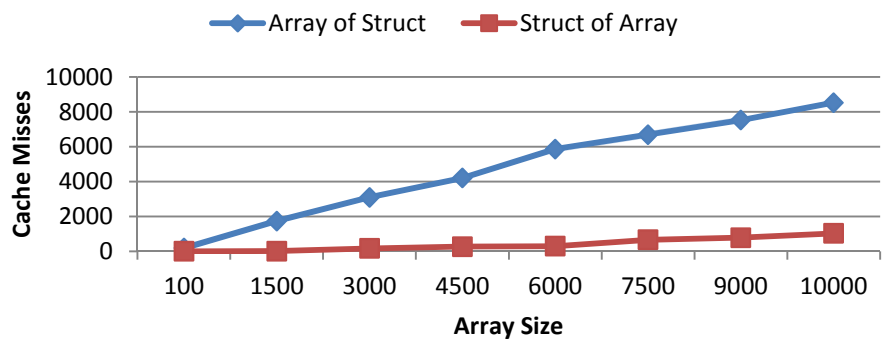
L1 Data Cache Misses



L2 Data Cache Misses

| Array Size | Array of Struct | Struct of Array |
|------------|-----------------|-----------------|
| 100 | 186 | 4 |
| 1500 | 1735 | 8 |
| 3000 | 3091 | 166 |
| 4500 | 4210 | 266 |
| 6000 | 5868 | 297 |
| 7500 | 6696 | 665 |
| 9000 | 7531 | 781 |
| 10000 | 8529 | 1026 |

L2 Data Cache Misses



Explanation:

- In case of Array of struct, there are 100 double variables (800 bytes) beside the variable x. Hence, when there is a compulsory miss for the first time, double variables are also fetched into the cache but never used. The subsequent value of x is available far apart and hence to access the value of x in the subsequent array index, we need to fetch the data from the memory into the cache. **There is no spatial locality.** Hence the number of misses keeps on increasing with the increase in the array size. Or in other words, the number of cache misses is more than the case of struct of array.
- In case of **Struct of Array**, the data which needs to be accessed are available next to each other (**Spatial locality**). Hence when the compulsory miss is encountered for the first time, the subsequent data is also fetched into cache, thereby avoiding future cache misses. Hence, the struct of array has less cache misses as compared to array of struct.
- Diagrammatically, both the data layout can be represented in the following form:

Data Layout 1 (Array of struct) –The data to be accessed are far apart. No spatial locality.



Data Layout 2 (Struct of array) – The data to be accessed are next to each other. Less misses due to spatial locality.

