

Operating Systems

MyThreads


User-level Threads

- Are used to allow programs to switch between multiple threads of execution without involving the operating system
- Use cooperative multitasking rather than preemptive multitasking
- Also consider co-routines
 - Subroutines are a special case that stop the calling routine until the called subroutine has completed (Knuth)

Jumping around in C

- It is often useful to quickly return to the top of a set of called functions when an error occurs

– Imagine recursive parsing



- **setjmp()** saves stack context for non-local *goto*
 - Non-local as a real goto is function scoped
 - Don't return from the function that called setjmp() or the context will be invalid

setjmp/longjmp

- Requires a global *jmp_buf jmpbuffer*;
 - Since it must be referenced in other functions
- if (**setjmp**(jmpbuffer) != 0) printf("returned\n");
 - That is, setjmp returns 0 when it is first called
- **longjmp**(jmpbuffer, 1);
 - Causes the setjmp to return a 2nd time, returning the value of the 2nd argument to longjmp

Jumping back

- What about the automatic variables on the stack? Variables in the registers?
 - It depends. The standards say it is “indeterminate”
 - Most implementations don’t try to roll back
- The **volatile** type qualifier in ANSI C tells the compiler to suppress optimizations that might remove apparently redundant memory accesses
 - `volatile int some_int;`

Think about stacks and variables


- If you allocate an automatic variable on the stack, that variable can only be used in that stack frame or “below”

```
int *  
open_file(void) {  
    int fd;  
  
    fd = open(. . .  
  
    return(&fd);  
}
```

Implementing User-Level Threads


- Requires a user-space (non system) mechanism to save and restore register sets
- How can this be done with what we’ve talked about?

Implementing Threads in the Kernel



A threads package managed by the kernel

Implementing Threads in User Space




A user-level threads package


One mechanism for User-Level Threads

- Stack management
- PC management
- Recall the *jmpbuffer*
 - It stores the PC, SP, BP or FP (among others)

Stacks for Kernel Threads



Stacks for User Level Threads



One thread approach

- Make a jmpbuffer
- Allocate a “stack” in the heap
- Modify jmpbuffer to point to your new stack
- Hold on to your hat...

Another way

- System V and later POSIX defined a more general way to manage context

```
typedef struct ucontext {
    struct ucontext *uc_link;
    sigset_t         uc_sigmask;
    stack_t          uc_stack;
        uc_stack.ss_sp...
        uc_stack.ss_size...
    mcontext_t       uc_mcontext;
    ...
} ucontext_t;
```

Context functions

- `int setcontext(const ucontext_t *ucp)`
 - Transfers control to the context in ucp
- `int getcontext(ucontext_t *ucp)`
 - Save the current context into ucp
- `void makecontext(ucontext_t *ucp, void *func(), int argc, ...)`
 - Create an alternate context with entry point func
- `int swapcontext(ucontext_t *oucp, ucontext_t *ucp)`
 - Transfer control to ucp and save current context into oucp


My Threads

- `void mt_init();`
- `void *mt_create(void (*func)(void *), void *arg);`
- `void mt_join(void *mt);`
- `void mt_joinall();`
- `void mt_exit();`
- `void mt_sleep(int seconds);`
- `void mt_yield();`
- `void *mt_self();`
- `void mt_kill(void *mt);`

My Threads Semaphores

- `typedef void *mt_sem;`
- `mt_sem mt_sem_create(int initial_value);`
- `void mt_sem_up(mt_sem sem);`
- `void mt_sem_down(mt_sem sem);`
- `void mt_sem_destroy(mt_sem sem);`
- `int mt_sem_getval(mt_sem sem);`

Stacks for User Level Threads



User-level threads

- Any thread call should invoke the user-level thread scheduler
- The first call to `mt_create` creates thread state and a stack
- Put the newly-created thread on the run queue
- The scheduler does any necessary housekeeping and chooses a runnable thread
 - In the beginning there is only one
- That thread runs until another thread routine is called
 - Then the scheduler is reentered

setjmp/longjmp

- Declare a global `jmp_buf jmpbuffer;`
- if (**setjmp**(jmpbuffer) != 0) printf("returned\n");
 - That is, `setjmp` returns 0 when it is first called
- **longjmp**(jmpbuffer, 1);
 - Causes the `setjmp` to return a 2nd time, returning the value of the 2nd argument to `longjmp`
- Use **volatile** for automatic variables

User-level threads

- Thread create must
 - malloc a stack for the new thread
 - Put the addresses of that stack in the jmpbuf
 - __jmpbuf[JB_BP] = newstack;
 - __jmpbuf[JP_SP] = newstack-1024;
 - Leave enough space for local variables

Context notes

- uc_link – when a context finishes (the entry function returns), the system swaps to the context pointed to by uc_link
- This could be a context that can clean up the thread state
- Scheduler can be a function or a distinct context

Joinall semantics

- Joinall will wait for all runnable threads
- If threads are blocked on semaphores or in sleep, joinall will return